


Action Model


Action Model

- How certain are we of our state after a given action?
- Action disperses the probability distribution



Modeling Action

- From odometry we have previous and current pose
 - $(x_{t-1}, y_{t-1}, \theta_{t-1})$
 - (x_t, y_t, θ_t)
 - $(\Delta x, \Delta y, \Delta \theta)$
- $\Delta s^2 = \Delta x^2 + \Delta y^2$
- $\alpha = \text{atan2}(\Delta y, \Delta x) - \theta_{t-1}$




- Model action as a rotation, translation and rotation:

$$u = [\alpha \quad \Delta s \quad \Delta \theta - \alpha]$$

Modeling Action Error

- Model the action as
 - $\text{Turn}(\alpha + \varepsilon_1)$
 - $\text{Travel}(\Delta s + \varepsilon_2)$
 - $\text{Turn}(\Delta\theta - \alpha + \varepsilon_3)$
- Model errors as Gaussian:
 - $\varepsilon_1 \sim \mathcal{N}(0, k_1 |\alpha|)$
 - $\varepsilon_2 \sim \mathcal{N}(0, k_2 |\Delta s|)$
 - $\varepsilon_3 \sim \mathcal{N}(0, k_1 |\Delta\theta - \alpha|)$
- Standard Deviation proportional to action magnitude



The Action Model $P(x_t | u_{t-1}, x_{t-1})$


- Given action: $u = [\alpha \quad \Delta s \quad \Delta\theta - \alpha]$

$$\begin{bmatrix} x_t \\ y_t \\ \theta_t \end{bmatrix} = \begin{bmatrix} x_{t-1} \\ y_{t-1} \\ \theta_{t-1} \end{bmatrix} + \begin{bmatrix} (\Delta s + \varepsilon_2) \cos(\theta_{t-1} + \alpha + \varepsilon_1) \\ (\Delta s + \varepsilon_2) \sin(\theta_{t-1} + \alpha + \varepsilon_1) \\ \Delta\theta + \varepsilon_1 + \varepsilon_3 \end{bmatrix}$$

- Where:
 - $\varepsilon_1 \sim \mathcal{N}(0, k_1 |\alpha|)$
 - $\varepsilon_2 \sim \mathcal{N}(0, k_2 |\Delta s|)$
 - $\varepsilon_3 \sim \mathcal{N}(0, k_1 |\Delta\theta - \alpha|)$

Actions Disperse the Distribution

- N particles is approximately a probability distribution
- The distribution disperses with action
- When tuning, look at your distribution and see if dispersion is reasonable given your experience with robot



Tuning the Action Model

- The noise in the action model captures non-systematic errors
- Calibrate odometry first to eliminate systematic errors.
- Can perform straight line experiments and rotation experiments to determine reasonable values of k_1 and k_2
- If dispersion is too small or too large, localization will fail.

Sample Odometry Action Model

```
1:  Algorithm sample_motion_model_odometry( $u_t, x_{t-1}$ ):  
2:       $\delta_{\text{rot1}} = \text{atan2}(\bar{y}' - \bar{y}, \bar{x}' - \bar{x}) - \bar{\theta}$   
3:       $\delta_{\text{trans}} = \sqrt{(\bar{x} - \bar{x}')^2 + (\bar{y} - \bar{y}')^2}$   
4:       $\delta_{\text{rot2}} = \bar{\theta}' - \bar{\theta} - \delta_{\text{rot1}}$   
  
5:       $\hat{\delta}_{\text{rot1}} = \delta_{\text{rot1}} - \text{sample}(\alpha_1 \delta_{\text{rot1}}^2 + \alpha_2 \delta_{\text{trans}}^2)$   
6:       $\hat{\delta}_{\text{trans}} = \delta_{\text{trans}} - \text{sample}(\alpha_3 \delta_{\text{trans}}^2 + \alpha_4 \delta_{\text{rot1}}^2 + \alpha_4 \delta_{\text{rot2}}^2)$   
7:       $\hat{\delta}_{\text{rot2}} = \delta_{\text{rot2}} - \text{sample}(\alpha_1 \delta_{\text{rot2}}^2 + \alpha_2 \delta_{\text{trans}}^2)$   
  
8:       $x' = x + \hat{\delta}_{\text{trans}} \cos(\theta + \hat{\delta}_{\text{rot1}})$   
9:       $y' = y + \hat{\delta}_{\text{trans}} \sin(\theta + \hat{\delta}_{\text{rot1}})$   
10:      $\theta' = \theta + \hat{\delta}_{\text{rot1}} + \hat{\delta}_{\text{rot2}}$   
  
11:     return  $x_t = (x', y', \theta')^T$ 
```

- Book uses 4 error coefficients α_1 to α_4
- α_1 and α_3 are same as k_1 and k_2 in our model
- α_2 and α_4 represent cross correlation
- You can use either model

Action Model Implementation


- Implemented in `action_model.cpp`
- Use `updateAction()` to check motion, calculate deltas and calculate standard deviations
- Use `applyAction()` to sample from normal distributions with previously calculated stdev and apply them to a particle.
- Lookup `std::normal_distribution` for how to sample from a normal distribution

Sampling from a Distribution


```
#include <random>
int main(){
    // initializing random seed
    std::random_device rd;
    // Mersenne twister PRNG, initialized with seed random device instance
    std::mt19937 gen(rd());
    int i;
    float sample;
    float mean = 0.0
    float stddev = 1.0
    // instance of class std::normal_distribution with specific mean and stddev
    std::normal_distribution<float> d(mean, stddev);
    for(i = 0; i < 1000; ++i)
    {
        // get random number with normal distribution using gen as random source
        sample = d(gen);
        do_something_with_this_value(sample);
    }
    return 0;
}
```

Sensor Model


Markov Localization



Prediction: apply action model



Sensor model



Correction: apply sensor model

Sensor Model $p(z_t|x_t)$

- Purpose: find the probability that a lidar scan matches the hypothesis pose (particle) given a map
- Give a score to each ray in the scan, sum for scan score
- Normalize score across particles to get particle weight
- Resample particle distribution based on weights before applying action model again.

Sensor Model




z_t, m





$p(z_t | x_t, m)$


Sensor Model


- Instead of calculating $p(z_t|x_t, m)$ for every possible pose in the map, only need to evaluate for each particle.
- Particle effectively represents hypothesis robot took the path from previous particle pose to new pose given by action model
- Sensor model will then test each ray in the laser scan and provide a fitness score based on the particles hypothesis
- Fitness scores normalized over all particles become particle weights
- You will decide how to assign fitness score based on accuracy/computational considerations
















Sensor Model: Beam Model

(a) Gaussian distribution p_{hit}




(b) Exponential distribution p_{short}





$$p(z_t^k | x_t, m) = \begin{pmatrix} z_{\text{hit}} \\ z_{\text{short}} \\ z_{\text{max}} \\ z_{\text{rand}} \end{pmatrix}^T \cdot \begin{pmatrix} p_{\text{hit}}(z_t^k | x_t, m) \\ p_{\text{short}}(z_t^k | x_t, m) \\ p_{\text{max}}(z_t^k | x_t, m) \\ p_{\text{rand}}(z_t^k | x_t, m) \end{pmatrix}$$

(c) Uniform distribution p_{max}



(d) Uniform distribution p_{rand}



Sensor Model: Beam Model

- The beam model looks at each ray endpoint casted on the grid and calculates the product of probabilities of the rays in the scan

```
1:   Algorithm beam_range_finder_model( $z_t, x_t, m$ ):  
2:        $q = 1$   
3:       for  $k = 1$  to  $K$  do  
4:           compute  $z_t^{k*}$  for the measurement  $z_t^k$  using ray casting  
5:            $p = z_{\text{hit}} \cdot p_{\text{hit}}(z_t^k \mid x_t, m) + z_{\text{short}} \cdot p_{\text{short}}(z_t^k \mid x_t, m)$   
6:                $+ z_{\text{max}} \cdot p_{\text{max}}(z_t^k \mid x_t, m) + z_{\text{rand}} \cdot p_{\text{rand}}(z_t^k \mid x_t, m)$   
7:            $q = q \cdot p$   
8:       return  $q$ 
```


Beam Model

- Measurement probabilities are too small to be meaningful

$$P(z_t|x_t) = \prod_{i=0}^N P(\text{ray}_i = d_i|x_t)$$

- Log odds can be accumulated without numerical underflow

$$\log P(z_t|x_t) = \sum_{i=0}^N \log P(\text{ray}_i = d_i|x_t)$$



Computing $\log P(\text{ray}_i = d_i | x_t)$

- One method is to cast the ray in the grid, and assign log odds to each case

- If ray_i terminates at the first obstacle

$$\log P(\text{ray}_i = d_i | x_t) = -4$$


- If ray_i terminates before an obstacle

$$\log P(\text{ray}_i = d_i | x_t) = -8$$

- If ray_i terminates after an obstacle

$$\log P(\text{ray}_i = d_i | x_t) = -12$$

- Take log odds into scan score




Sensor Model: Likelihood Field

- Overcomes lack-of-smoothness and computational limitations of Sensor Beam Model
- Idea: Instead of ray casting (computationally expensive) just check the end-point.
- The likelihood $p(z|x_t, m)$ is given by a Gaussian distribution

$$\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{d^2}{2\sigma^2}\right)$$

Where d is the distance to the nearest obstacle and σ is the std dev.




Sensor Model: Likelihood Field

```
1:  Algorithm likelihood_field_range_finder_model( $z_t, x_t, m$ ):  
2:       $q = 1$   
3:      for all  $k$  do  
4:          if  $z_t^k \neq z_{\max}$   
5:               $x_{z_t^k} = x + x_{k,\text{sens}} \cos \theta - y_{k,\text{sens}} \sin \theta + z_t^k \cos(\theta + \theta_{k,\text{sens}})$   
6:               $y_{z_t^k} = y + y_{k,\text{sens}} \cos \theta + x_{k,\text{sens}} \sin \theta + z_t^k \sin(\theta + \theta_{k,\text{sens}})$   
7:               $dist = \min_{x', y'} \left\{ \sqrt{(x_{z_t^k} - x')^2 + (y_{z_t^k} - y')^2} \mid \langle x', y' \rangle \text{ occupied in } m \right\}$   
8:               $q = q \cdot \left( z_{\text{hit}} \cdot \mathbf{prob}(dist, \sigma_{\text{hit}}) + \frac{z_{\text{random}}}{z_{\max}} \right)$   
9:      return  $q$ 
```

Simplified Likelihood Field Model

- Look at endpoints (like likelihood field)
- Take log odds of cell into scan score if positive
- If it is not a hit, check cell before and after along ray and take fraction of log odds into scan score



SLAM Implementation

- Initialize Particle Filter

- Loop:

```
updateLocalization();  
    updateFilter(odometry, lidar, map);  
        resamplePosteriorDistribution(); // resampling using weights  
        computeProposalDistribution(); // apply action model  
        computeNormalizedPosterior(); // apply sensor model  
        estimatePosteriorPose(); // i.e. weighted average  
updateMap(); //update with posterior pose
```

Moving Laser Scan in Sensor Model

```
double SensorModel::likelihood(...){  
  
    double scanScore = 0.0;  
    MovingLaserScan movingScan(scan, sample.parent_pose, sample.pose);  
  
    for(const auto& ray : movingScan)  
    {  
        rayScore=scoreRay(ray);  
        scanScore += rayScore;  
    }  
    return scanScore;  
}
```

Resampling Particles

- We want to resample the particle distribution at each iteration
- Sample based on weight of each particle
- Sampling is done *with replacement* – we can potentially sample high weighted particles multiple times
- Repetitive resampling amplifies the variance (not good)
- Should not resample if the robot does not move
- If we randomly chose weighted particles to resample independent of each other we risk losing diversity of the particles

Low Variance Resampling

Algorithm Low_variance_sampler($\mathcal{X}_t, \mathcal{W}_t$):

$\bar{\mathcal{X}}_t = \emptyset$

$r = \text{rand}(0; M^{-1})$

$c = w_t^{[1]}$

$i = 1$

for $m = 1$ *to* M *do*

$U = r + (m - 1) \cdot M^{-1}$

while $U > c$

$i = i + 1$

$c = c + w_t^{[i]}$


endwhile

add $x_t^{[i]}$ *to* $\bar{\mathcal{X}}_t$

endfor

return $\bar{\mathcal{X}}_t$

- For resampling M particles, Choose a random number r
- Select particles by repeatedly adding M^{-1} to r and choosing the corresponding particle
- c accumulates weights to handle resampling higher weighted particles more often



Losing Diversity

- Loss of diversity can be caused by sampling a discrete distribution
- Solution: “regularization”
 - Consider particles represent a continuous distribution
 - Sample from the continuous density

- Example: Given 1D particles: $\{x^{(1)}, x^{(2)}, \dots, x^{(K)}\}$

- Sample from the density:
$$p(x) = \sum_{k=1}^K \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-x^{(k)})^2}{\sigma^2}}$$

Particle Deprivation

- Happens when there are no particles in the vicinity of the correct state
- Occurs as the result of the variance in random sampling. An unlucky series of random numbers can wipe out all particles near the true state.
 - This has non-zero probability to happen at each time so it will happen eventually
- Popular solution: add a small number of randomly generated particles when resampling
 - Advantage: reduces particle deprivation, is simple
 - Disadvantage: incorrect posterior estimate even in the limit of infinitely many particles

Posterior Pose Estimate

- Particle representation of distribution may not contain true pose, especially with few particles.
- You will need to come up with a robust method to estimate the posterior pose to feed into the mapping step.
- Many options:
 - Weighted average of poses?
 - Average of “good” poses (i.e. top 10% of weights)?
 - K-means clustering?
 - Depends on the variance of distribution?

Averaging Angles

- Arithmetic mean is not appropriate for mean of angles
- Find mean in Cartesian coordinates and convert back with atan2.

$$\bar{\theta} = \text{atan2}\left(\sum_{i=0}^N w_i \sin(\theta_i), \sum_{i=0}^N w_i \cos(\theta_i)\right)$$

Putting it together

- Mapping
 - `slam -mapping-only`
 - `botgui`
 - `log-player-gui <log-file>` (turn off SLAM_MAP & SLAM_PARTICLES)
- Action Model
 - `slam -action-only -localization-only <map_file>`
 - `botgui`
 - `log-player-gui <log-file>` (turn off SLAM_MAP & SLAM_POSE & SLAM_PARTICLES)
- Sensor Model
 - `slam -localization-only <map_file>`
 - `botgui`
 - `log-player-gui <log-file>` (turn off SLAM_MAP & SLAM_POSE & SLAM_PARTICLES)
- Full Slam
 - `slam`
 - `botgui`
 - `log-player-gui <log-file>` (turn off SLAM_MAP & SLAM_POSE & SLAM_PARTICLES)

A problem with the logs

- RPLidar driver changed from W22 to F22
- Before – LIDAR scan clockwise
- After – LIDAR scan counter clockwise
- Solution:
 - Add a command line arg for direction of LIDAR
 - If -CW use $2\pi - \theta_{ray}$
 - If -CCW use θ_{ray}
 - Old logs will be CW, but new data will be CCW